

## Part 1: Building Language Models

Firstly, we loaded the data using the `load_jsonl` function.

After loading the data, we filtered the sentences based on the protocol type (e.g., **committee** and **plenary**) using the `get_sentences_by_type` function.

Then to prepare the data for training the Trigram model, we added start tokens (`<s_0>` and `<s_1>`) at the beginning of each sentence using the `add_start_tokens` function.

After that, We implemented the **Trigram Language Model** using the `Trigram_LM` class. This involved the following steps:

- Calculating **unigram counts**: The frequency of each individual token in the corpus.
- Calculating **bigram counts**: The frequency of consecutive token pairs.
- Calculating **trigram counts**: The frequency of token triplets.
- Storing the vocabulary size and total token count for smoothing purposes.

---

### Function: `calculate_prob_of_sentence`

---

Within the `Trigram_LM` class, we implemented the `calculate_prob_of_sentence` function. This function computes the log probability of a given sentence using:

- **Unigram probabilities.**
- **Bigram probabilities.**
- **Trigram probabilities.** To handle unseen words and ensure non-zero probabilities, we applied **Laplace smoothing**. Probabilities were combined using **linear interpolation** with predefined weights ( $\lambda_1, \lambda_2, \lambda_3$ ). Higher weight was given to the trigram because it captures the most context-specific information. Bigrams and unigrams received lower but non-zero weights to handle sparsity and provide a fallback.  
We decided to choose :  $\lambda_1=0.1, \lambda_2=0.2, \lambda_3=0.6 \rightarrow \lambda_1 + \lambda_2 + \lambda_3 = 1$

---

### Function: `generate_next_token`

---

Within the `Trigram_LM` class, we implemented the `generate_next_token` function.

This function was implemented to predict the most likely next token given a two-token context:

- Calculates trigram, bigram, and unigram probabilities using MLE.
- Applies linear interpolation to combine these probabilities.

- Returns the token with the highest log probability and its corresponding probability value.

In cases where no valid token was found, the model defaults to the most frequent unigram as a fallback mechanism.

## Part 2: Collocations

We implemented the `get_k_n_t_collocations` function within the `Trigram_LM` class, which operates by extracting n-grams, filtering them based on a threshold `t`, and ranking the top `k` n-grams using **Frequency** or **TF-IDF** scoring methods.

This function supports two ranking methods:

1. Frequency-Based Ranking

- The remaining n-grams are sorted in descending order of their occurrence counts.
- The top `k` most frequent n-grams are returned

2. TF-IDF-Based Ranking

To rank the n-grams using TF-IDF, the function calls `calculate_tfidf`, which calculates the scores based on the TF-IDF formula provided in the task.

This function works as follows:

- It computes the **TF** of each n-gram, which measures how often the n-gram appears in a sentence relative to all n-grams in that sentence.
- It calculates the **DF** to determine how many sentences contain the n-gram.
- Using TF and DF, the function computes the **TF-IDF score** for each n-gram, assigning higher importance to n-grams that are frequent in individual sentences but not common across all sentences.

As required, we applied a **threshold `t`** to filter out n-grams that appear fewer than `t` times:

```
# frequency threshold `t`
filtered_ngrams = {ngram: count for ngram, count in ngram_counts.items() if count >= t}
```

At the end, the top `k` n-grams with the highest scores are printed to the `kneset_collocations.txt` file as requested.

## **Part 3: Applying the Language Models**

### **Step 1:**

We implemented the `mask_tokens_in_sentences` function. This function processes a list of sentences and replaces a specified percentage of tokens  $x$  in each sentence with the special placeholder token [ \* ].

The function operates as follows:

#### **1. Calculate Number of Tokens to Mask**

- For each sentence, the number of tokens to be masked is determined by multiplying the length of the sentence by  $x$ .
- If the result is a decimal, it is rounded up or down based on standard rounding rules.
- At least **one token** is always masked, even if  $x$  is very small.

#### **2. Random Selection of Tokens**

- Random indices are selected within the sentence to determine which tokens will be replaced.
- The `random.sample` method ensures no duplicate indices are chosen.

#### **3. Token Replacement**

- The selected tokens are replaced with the placeholder [ \* ] in a copy of the sentence to avoid modifying the original input.

#### **4. Store Masked Sentences**

- The function returns a list of masked sentences, where each sentence includes the same structure as the original but with the masked tokens replaced.

## **Step 2:**

We implemented the `sample_and_save_sentences` function. This function processes a list of sentences, filters them, and saves both the original and masked versions to files.

this function operates as follows:

### **1. Sentence Filtering**

- The function selects only sentences that contain **at least 5 tokens**.

### **2. Sentence Sampling**

- From the filtered sentences, the function randomly samples `num_samples` sentences.

### **3. Token Masking**

- For each sampled sentence, `mask_tokens_in_sentences` is called to replace **x%** of the tokens with the placeholder [ \* ].

### **4. Saving Results to Files**

- The original sentences are saved to a file named `original_sampled_sents.txt`
- The masked sentences are saved to a file named `masked_sampled_sents.txt` as the order of sentences in the in the original file.

### **Step 3:**

We implemented the `restore_and_evaluate_sentences` function.

this function operates as follows:

1. **Load Original and Masked Sentences**
  - The function reads the **original sentences** and the corresponding **masked sentences** from the input files.
2. **Restore Missing Tokens**
  - For each masked sentence, tokens marked as [ \* ] are restored using the **plenary model**.
  - The restoration process works by predicting the masked token based on its **context**:
    - $w_{k-2}$  (two tokens before the mask)
    - $w_{k-1}$  (one token before the mask)
  - If no prior context exists, special tokens <s\_0> and <s\_1> are used.
3. **Track Generated Tokens**
  - The restored tokens are collected and stored in a comma-separated list.
4. **Evaluate the Restored Sentences**
  - The probabilities of the restored sentences are calculated using both:
    - The **plenary model** (trained on plenary corpus).
    - The **committee model** (trained on committee corpus).
5. **Save the Results**
  - The results for each sentence are written to the `sampled_sents_results.txt` file.

## Step 4:

We implemented the `calculate_perplexity` function to measure how well the plenary language model predicts the masked tokens. Perplexity serves as a quantitative measure of the model's uncertainty when restoring the tokens, where a lower perplexity indicates better predictions.

**PERPLEXITY**

- The probability assigned by the language model to the test set, normalized by the size of the test set:
$$\text{perplexity}(x_{1..n}, m) = 2^{H(x_{1..n}, m)} = m(x_{1..n})^{-\frac{1}{n}}$$
- We use the following notation:
$$PP(w_{1..N}) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}} = \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}$$
- By the chain rule, this is:
$$\sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 w_2 \dots w_{i-1})}}$$
- Which for bigrams gives:
$$\sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_{i-1})}}$$

Shuly Wintner (University of Haifa)    Natural Language Processing    ©Copyrighted material    287 / 728

We used this definition of perplexity to evaluate the performance of the plenary language model. Specifically:

### 1. By the Chain Rule:

The probability of a sentence can be broken down into smaller components, where each token's probability depends on its preceding tokens. This allows us to approximate the overall probability of a sentence using n-grams.

### 2. Simplification for Bigrams:

For simplicity, we considered each token's probability based on its immediate previous token. This approach makes it computationally feasible to evaluate perplexity.

### 3. Perplexity Calculation:

Perplexity is calculated as a measure of the model's uncertainty in predicting the masked tokens. It considers the average probability of the restored tokens within their context:

- Higher probabilities mean the model is more confident in its predictions.
- Lower perplexity scores indicate better performance because the model is less "surprised" by the restored tokens.

### 4. Interpolation of Probabilities:

To make the predictions more robust, we combined three probabilities:

- **Unigram probability:** Based on the frequency of a single token.

- **Bigram probability:** Based on the frequency of a token given the previous token.
- **Trigram probability:** Based on the frequency of a token given the two previous tokens.

We assigned weights to each of these components:

- 10% to unigram, 30% to bigram, and 60% to trigram probabilities.

#### 5. Output:

- The average perplexity was computed for all sentences, considering only the masked positions.
- The final result was saved to the file **perplexity\_result.txt**.
- The average perplexity for the restored sentences is **17,243.48**.

The average perplexity for the restored sentences, with the current lambda values ( $\lambda_1 = 0.1, \lambda_2 = 0.3, \lambda_3 = 0.6$ ), is **17,243.48**.

If the lambda values are modified, the perplexity score changes to reflect the new weighting of unigram, bigram, and trigram probabilities.